RESEARCH NOTE / 16 OCTOBER 2026

A cheaper token is not always a cheaper task.

Three reproducible examples show why API price comparisons need a tokenizer, an effective date and an output budget. These are accounting scenarios, not model-quality rankings.

01 TOKENIZER

Same input text can become more billable tokens.

Hypothetical: 1M old input tokens; 33.3% lower unit price, 13.3% lower cost.

02 EFFECTIVE DATE

Gemini 4 Argon: the effective date changes the cost.

1M input + 0.2M billable output tokens, aggregated across requests.

03 OUTPUT MIX

Gemini 3.5 Flash-Lite: output can dominate.

Same aggregate mix as case 02; output contributes $0.50, or 62.5%.

Bar scales differ between cases; compare lengths only within each case.

EDITORIAL TAKEAWAY

Ask what stays constant. Equal token counts do not guarantee equal text; current rates need not survive a budget period; and input-only prices omit output. Report the exact accounting basis alongside any savings claim.

METHOD & SOURCE CHECKS

Reproduce the comparison.

Standard paid API rates, USD per million tokens; observed 16 October 2026.

Model	Input	Output	Effective period
Sonnet 4.6 [1]	3.00	15.00	Current; legacy
Sonnet 5.5 [1]	2.00	10.00	Current
Gemini 4 Argon [2]	2.00	10.00	To 30 Nov 2026
Gemini 4 Argon [2]	4.00	20.00	From 1 Dec; scheduled
Gemini 3.5 Flash-Lite [2]	0.30	2.50	Current

WHAT WAS CALCULATED

Cost = input millions x input rate + billable output millions x output rate. Quantities are aggregate across requests, never a single oversized request.

Case 01 uses an assumed 1.30 input-token multiplier, based on Anthropic's approximate newer-tokenizer comparison. $2 x 1.30 = $2.60 versus $3.00: 13.3% lower. This is a sensitivity example, not a measured corpus result.

Case 02 holds billable volumes fixed: $2 + (0.2 x $10) = $4; scheduled $4 + (0.2 x $20) = $8. Case 03: $0.30 + (0.2 x $2.50) = $0.80. Output share: $0.50 / $0.80 = 62.5%.

VERIFICATION & LIMITS

Rates were checked against provider pricing tables; model pages confirmed the selected current models and Sonnet 4.6's legacy status. A second AI agent checked the source interpretation and arithmetic. The accompanying CSVs and Python script reproduce every plotted value.

No live bills, tokenization, latency or quality were measured. Google output counts include thinking tokens. Caching, batch/flex/priority rates, regional premiums, tool charges, taxes and negotiated discounts are excluded. Future prices are the published schedule observed today, not a guarantee.

PRIMARY SOURCES

[1] Anthropic model pricing and tokenizer note

https://platform.claude.com/docs/en/about-claude/pricing

[2] Google Gemini Developer API pricing

https://ai.google.dev/gemini-api/docs/pricing

Verification session: 15:36-15:42 UTC, 16 October 2026. Provenance in JSON.